

Computer Network Defense

we begin with the triad

- patch updates

 - #1 most overlooked security technique

 - as vulnerabilities are found, patches are released

 - how long (on average) do you think a vulnerability exists before it is discovered?

 - 342-ish days (uhhh, yeah, almost one year!)

 - patches may make you safe from *most* attacks

 - just not 0-day attacks

- malware protection

 - no, you are not invincible!

 - no matter what you think

 - malware?

 - viruses

 - worms

 - bacteria

 - trojans

 - rootkits

 - spyware (e.g., sniffers, keyloggers, etc)

 - adware (e.g., annoying popups, spam, phishing, etc)

 - defense?

 - anti-virus

 - anti-spam/anti-adware

 - e.g., spybot s&d, avast, ad-aware, avg, comodo, mcafee, norton, clamav

 - anti-malware?

 - hash detection, basically

- firewall

 - take care of what's on your system

 - two philosophies

 - block based on port (which usually ties to services/protocols)

 - any application attempting connection on a port is blocked

 - this is the Linux way

 - block based on application

 - a single application is blocked

 - this is the Windows way

 - usually means having to interact more with the firewall

 - which is usually annoying

 - stateful vs. stateless

 - do we treat each packet uniquely (no past memory)? → stateless

 - or do we use the past to infer something about the now? → stateful

 - h/w (e.g., router, dedicated Linux firewall) vs. s/w (e.g., zone alarm, windows firewall, etc)

 - how cool would it be to be able to “program” hardware networking equipment?

 - we have some FPGA-based devices to let us create any device—and program it!

 - you can also setup a Linux box as the front facing thing in your network

setup iptables/netfilter properly

iptables: the Linux interface to netfilter

netfilter: the core Linux firewall

any downloaded firewall simply makes interacting with iptables more “friendly”

e.g., firestarter, ufw, etc

defensive operations

what can we do to “protect” ourselves?

one option is to encapsulate our services/OS

e.g., virtualization (hypervisor and virtual machines)

e.g., virtualbox, vmware, xen, proxmox, etc

e.g., chroot jails (**see the relevant document on the class web site**)

defense in depth

don't depend on a single mechanism for protection

layered approach (multiple layers of defense) – like an onion in your network

idea: use several varying methods

e.g. anti-virus on router, anti-virus on firewall, anti-virus on downstream machines

sometimes, we delay rather than prevent: yield space in order to buy time

so it should prevent security breaches while giving time to respond

can we draw parallel to DFS in data structures?

defense in breadth

there are many attack vectors (i.e., just having a firewall won't guarantee security)

we must try to cover all attack vectors

e.g., email security and messaging security and anti-virus and spyware, and ...

can we draw parallel to BFS in data structures?

IDS/IPS

how can we detect intrusions?

how can we detect attackers?

could we protect/prevent in addition to detect?

these usually inspect packets (sometimes deeply)

tcp wrappers (rules)

maybe we can think about this being like a filter for tcp packets

we can scan, log, anonymize, etc

and maybe we could detect/protect/prevent via tcp wrappers

PDR³ (or should it be PDRER?)

prevent

we're a “pill” society

we prefer to take care of the symptoms, not the cause

and that's often a bad idea (but a money-making one!)

better idea: identify the cause and prevent the problem from occurring again

but that takes work (effort) – that's why we're a symptom-based society

we're lazy people, when you think about it

so best case is to prevent security breaches and vulnerability exploits

but that's not always possible, particularly in cyberspace

detect

if we can't prevent, we must find out when we have a problem

so use an ids, ips, idps

and also firewall, patches, anti-virus (i.e., the triad)

respond

if we detect, we can't just let something bad happen

what to do, what to do?!

how proactive can we be?

do we just secure our system and repair?

then prevent the perpetrator from doing it again (how?)

can we “engage?”

can we find out who did this and where they live?

recover

if our system was compromised, we may need to recover

how might we do this?

or might we endure instead of recover? or both?

restore

maybe our system is irrecoverable

so we take this as a learning experience

we restore from some previous backup

then we look at how to prevent this from happening again

and we loop back to the beginning...

avoid?

how the hell do we do this?

threat avoidance

threats simply don't matter

we don't care about detection, mitigation, prevention, attribution

we have an invisibility cloak

e.g., beaconing malware, unauthorized network users/apps, port knocking

honeypots prove useful

they have no production value

they lure attackers

we want to know what they do, what they use, how they do it

quasi-honeypot

make it more “useful?”

a part has production value

traffic patterns are like a production machine

attackers are less likely to detect that it's a honeypot

honeynets

interesting and relevant types of attacks that we may have to defend against

ddos (the holy grail)

dos: denial of service attack

attempt to make computer resources unavailable

ddos: originates from multiple systems

how?

consume computer resources (bandwidth, cpu, disk space)

disrupt configuration information (e.g. routing information)

disrupt state information (e.g. reset tcp sessions)

disrupt physical network components

obstruct communication

smurf attack (ping flood)

- generate a lot of network traffic on a network

- by flooding the target system with spoofed ping messages to broadcast addresses

syn flood

- SYN, SYN-ACK, ACK, hang up

- half-open connections may take up resources on the client

ping of death

- normal ping packet size is 56 or 84 bytes

- sending one that is larger than max ip packet size (65,535 bytes) could cause a crash (old)

pdos: permanent dos

- phlashing (illegitimate flashing of hardware → bricks the device)

application level floods

- irc floods

- buffer overflows

- banana attack: redirect outgoing messages back to sender

degradation of service

- many zombies mount temporary dos

- harder to detect

some are unintentional (e.g., google news on the day of michael jackson's death)

backscatter

- some attackers spook source ip

- you respond as usual

- those response packets are backscatter

- imagine if I spoofed millions of packets with your address as the source?

botnets

- a bunch of zombies!

- software agents that run autonomously and automatically

- mostly interpreted to be malicious; but can be legitimate (e.g., SETI)

- compromised via

- drive-by-downloads (RTFM!)

- awareness is important (in everything actually)

- browser exploits (IE6)

- worms

- trojans

- backdoors

- bot herder/master establishes C3

- often takes place on irc server

- usually runs hidden in a covert channel

- Dutch police once found a 1.5 million node botnet!

- they are now larger!

- used in many ways and typically auctioned to highest bidder; for:

- spam, ddos, click fraud, adware, spyware, etc

script kiddies

- those who use scripts or programs developed by others to attack computer systems
 - but they really don't know anything more than that
 - no knowledge of the underlying concepts
 - they're just annoying
- most “hackers” are actually script kiddies
- tools they like to use
 - winnuke (dos)
 - back orifice (remote system administration)
 - netbus (remote system administration)
 - sub7 (remote system administration)
 - netbus backwards and then substitute 7 with ten
 - 1/1/2010: hacker took them down (still down and closed forever)
 - metasploit (os computer security project)
 - prorat (backdoor Trojan)
 - and more...

network telescope (darknet, internet motion sensor, black hole)

- used to take a look at the unused part of the Internet
- all traffic to these addresses is suspicious
- some ip addresses are carved out from the Internet addressable space
 - 127.0.0.0 through 127.255.255.255 (loopback)
 - 192.168.0.0 through 192.168.255.255
 - 10.0.0.0 through 10.255.255.255
 - 172.16.0.0 through 172.31.255.255
 - and more... (see https://en.wikipedia.org/wiki/Reserved_IP_addresses)

darknet

- a portion of routed ips with no active services/servers
- seemingly nothing there
- but there is at least one server that accepts all packets (a packet vacuum)
- packets do arrive there, but by mistake or misconfiguration
- but it's primarily malware that goes there
- useful for analysis
- also nefarious servers

offensive operations

- sometimes, the best defense is knowledge of the offensive tactics
- so many of the things here are often employed defensively
 - let me see what I can gather from my own networks and systems from an offensive side
 - so that I can build a better defensive side

reconnaissance and footprinting

- useful to see if we might want to gain access to a system we don't have access to
- recon: what's there? what systems exist?
- footprinting: what specific things can we gather about those systems?
- we might want to know a few things about a system:
 - what OS it runs
 - what hardware it has
 - what servers are running and on what ports
 - including versions of all of these (some may have known vulnerabilities!)

forums and testimonials are a good resource
funny how many it techs post their problems (help!) and system specifics online
we give out too much information
do you facebook? myspace?
once we have this, we can head over to many exploits databases
anti-virus/anti-spam/firewall testimonials tell us what people are using
if exploits exist, we might find a way in
many malwares try to shut off protection software like anti-virus, etc

recon tools

nmap: security scanner for network exploration and security audits
strobe: essentially an fast and efficient nmap (on steroids)
nemesis: packet crafter and sender
python-scapy: packet swiss army knife
netcat: tcp swiss army knife
telnet: not as good as netcat

recon/footprinting tactics

port scanning

probes remote host for open ports
used to verify security policies and identify running services
portscan: scan for listening ports
portsweep: scan multiple hosts for a specific port
some worm may portsweep many hosts for a single port (vulnerability)

port status

open/accepted: something is listening
closed/denied/not listening: connection is denied
filtered/dropped/blocked: no reply (firewall?)

tcp scanning

use OS network functions
in nmap, called a connect scan
on connect, handshake performed and connection closed
no special privileges required
no low level control

syn scanning

uses raw ip packets and monitors for responses
known as half-open scanning because never actually opens a full TCP connection
port scanner generates a SYN packet
if target port is open, host responds with SYN-ACK
port scanner responds with RST and closes connection before handshake
we can get many details this way
target service never actually receives a connection
usually requires privileges

udp scanning

udp is a connectionless protocol
response comes only if a port is closed
so absence of response implies port is open
most scanners use this method
firewalls can fool scanner

ack scanning

- does not determine whether a port is open/closed
- instead, if it is filtered
- useful to probe for firewalls and its rulesets

network sniffing (particularly under the same subnet) – “sniffer”

- packet analysis
- intercepts/logs network traffic (packets)
- we can then decode/analyze these packets
- uses:
 - analyze network problems
 - detect network intrusion attempts
 - gain info for possible network intrusion
 - monitor network usage
 - gather/report network stats
 - filter content from traffic
 - spy on users/collect sensitive information
 - reverse engineer proprietary protocols
 - debug client/server communication
 - debug network protocols

tools

- carnivore
 - FBI's version
 - designed to monitor email and electronic communication
- tcpdump, wireshark (formerly ethereal)
- python-scapy (wrap your own sniffer around it)
- cain and abel
 - mainly a password recovery tool
 - but can sniff passwords transmitted through packets
 - exhaustive methods to “recover” passwords

arp spoofing

- arp = address resolution protocol
- can be used to poison (arp poisoning)
- man-in-the-middle
 - can stop traffic
 - can modify traffic
- can only be used on networks that make use of arp
- tools
 - ettercap

more about nmap

- it's a network mapper
- it creates a “map” of the network
- features:
 - host discovery
 - port scanning
 - version detection
 - OS detection
- used to:

perform security audits (identify connections, identify unexpected new servers)
identify open ports
get a network inventory